

NLP Multiple Choice Questions

Section 1 - 5

1. You are building a chatbot that answers customer questions using AI. Which field of AI primarily focuses on making this possible?
 - A. Computer Vision
 - B. Natural Language Processing
 - C. Reinforcement Learning
 - D. Robotics

2. In one-hot encoding, how many dimensions would the vector have if the vocabulary size is 7?
 - A. 3
 - B. 7
 - C. Number of unique letters
 - D. Depends on sentence length

3. Given a dataset with repeated words like "I, dogs, agree, agree", which representation would best capture the repetition?
 - A. One-hot encoding
 - B. Bag of Words
 - C. Word Embeddings
 - D. ASCII encoding

4. Which representation method best captures semantic meaning?
 - A. One-hot encoding
 - B. Bag of Words
 - C. Embedding vectors
 - D. ASCII encoding

5. Which representation requires training (unsupervised learning)?
 - A. One-hot encoding
 - B. Bag of Words
 - C. Word Embeddings
 - D. Character-level counts

6. When visualizing pretrained embeddings using PCA for the following words: (man, woman, male, female, king, queen, prince), what are you likely to see?
 - A. Word frequency distribution
 - B. Clusters of semantically related words
 - C. Binary patterns of text
 - D. Randomly distributed words

7. From the following, which is the most likely representation for the words "fast" and "quick"?

- A. fast: [0.88, 0.45, 0.12, 0.91] / quick: [0.85, 0.41, 0.15, 0.89]
- B. fast: [0.88, 0.45, 0.12, 0.91] / quick: [-0.88, -0.45, -0.12, -0.91]
- C. fast: [1.0, 0.0, 0.0, 0.0] / quick: [0.0, 1.0, 0.0, 0.0]
- D. fast: [0.10, 0.95, -0.30, 0.05] / quick: [0.85, -0.15, 0.70, -0.50]

8. You want to find all occurrences of the word "cat" in a text. Which regex should you use?

- A. cat
- B. Cat
- C. [cat]
- D. (cat)

9. What does the regex \d match?

- A. Any lowercase letter
- B. Any digit
- C. Any punctuation
- D. A space

10. To find a sentence ending with "dog", which pattern should you use?

- A. ^dog
- B. dog\$
- C. .*dog
- D. \b dog\b

11. To remove all punctuation marks, which regex can be used with re.sub()?

- A. r"\s"
- B. r"[^\w\s]"
- C. r"\d"
- D. r"."

12. Which library is known for its fast pipelines and pretrained models?

- A. NLTK
- B. spaCy
- C. Stanza
- D. TextBlob

13. If your goal is research and experimentation with tokenizers, corpora, and linguistic algorithms, which should you use?

- A. NLTK
- B. spaCy
- C. Stanza
- D. TextBlob

14. Which framework is most commonly associated with Transformer models in NLP?

- A. NLTK
- B. spaCy
- C. Hugging Face Transformers
- D. TextBlob

15. Which statement best describes CUDA?

- A. It's used for CPU parallelization
- B. It enables GPU-based general computation
- C. It is a visualization tool
- D. It is a text-processing API

16. Given the dataset below, which class is the Naive Bayes model most likely to predict for the sentence "amazing story and acting"?

ID	Label	Text
1	Positive	love movie amazing acting
2	Positive	great story love characters
3	Positive	amazing and wonderful soundtrack
4	Negative	boring movie slow
5	Negative	terrible acting waste of time
6	Negative	boring story and bad script

- A. Positive
- B. Negative
- C. Both equally
- D. Neutral

17. What is the main goal of text classification?

- A. Generating new text
- B. Translating text into another language
- C. Assigning labels or categories to text
- D. Extracting entities from text

18. Why do modern NLP classifiers often use embeddings instead of one-hot encoding?

- A. Embeddings are slower but simpler
- B. Embeddings capture semantic meaning
- C. One-hot encoding is more memory-efficient
- D. Embeddings remove context

19. What does Laplace Smoothing solve in Naive Bayes classification?

- A. It speeds up computation
- B. It prevents zero probabilities for unseen words
- C. It increases model complexity
- D. It improves tokenization

20. What is the main advantage of neural networks over Naive Bayes in text classification?

- A. They train faster
- B. They capture nonlinear patterns and contextual meaning
- C. They require less data
- D. They use fewer parameters

21. What is the main goal of text preprocessing in NLP?

- A. To visualize text data
- B. To clean and prepare raw text for modeling
- C. To train neural networks
- D. To remove all words from the text

22. Which step in preprocessing helps convert different forms of a word to a single root form?

- A. Tokenization
- B. Normalization
- C. Stopword Removal
- D. Spell Checking

23. Which library is widely used for spell checking and is considered *superior* for fast correction?

- A. SymSpell
- B. TextBlob
- C. SpaCy
- D. NLTK

24. Which of the following would lemmatization convert correctly, but stemming might not?

- A. playing → play
- B. mice → mouse
- C. studying → study
- D. worked → work

25. What does tokenization do?

- A. Converts text into numerical form called tokens
- B. Splits text into smaller units
- C. Removes stopwords
- D. Checks spelling

26. Dataset sample: "BUY CHEAP TOYS NOW!!!" and "Meeting tomorrow at 10 am."

Which cleaning step will best help you remove noise from this dataset?

- A. Lowercasing
- B. Tokenization
- C. Normalization
- D. Spellchecking