

NLP -Section 9

Self-Attention + Transformers

Self-Attention

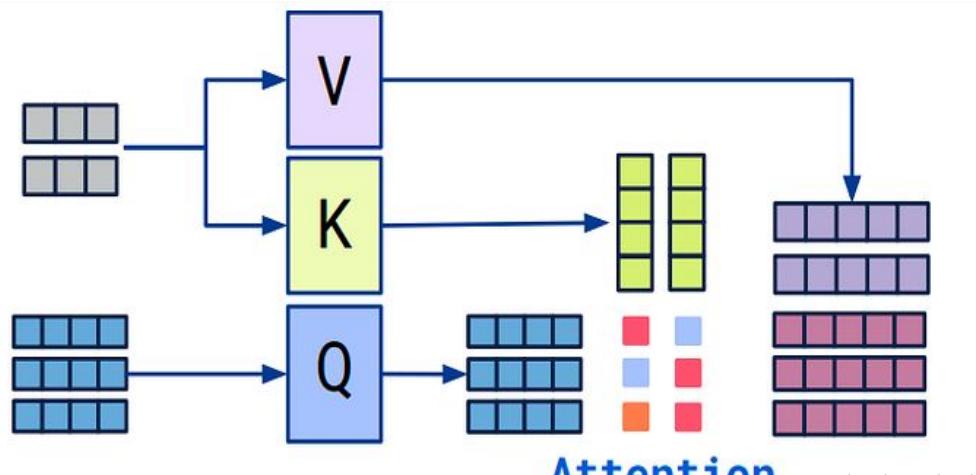
- Same idea as attention but instead of reflecting the attention on the output, it reflects on itself.
- Creates context embeddings
- Lacks positionality

Self Attention - Matrices

- We create Q , K , V for each token.
- Q (query): A learned projection of the input embedding used to compute attention scores toward all other tokens.
- K (key): A learned projection that interacts with queries so that the model can decide how much attention a token should receive.
- V (value): A learned projection that carries the content to be aggregated based on the attention weights.
- Q , K , and V have # of columns equal to the number of tokens each column representing a token (by small letter and index ex: q_1, v_3).
- Calculated through Wq^*X , Wk^*X , Wv^*X

Self-Attention - Calculation

- The most straightforward method:
 1. Calculate the score by $q \cdot k$ (the dot product) to get a single value
 2. Weight the values around zero by dividing by root d then applying softmax
 3. Multiply by v to get the output context representation z



$$\text{softmax} \left(\frac{Q \times K^T}{\sqrt{d_k}} \right) = Z$$

Below the equation, the components are labeled: Q , K^T , V , and Z .

The self-attention calculation in matrix form

Transformers Architecture - Positional Encoding

- Positional Encoding can be learned (new architectures) or sinusoidal (original transformer).
- Sinusoidal where:

$$PE(pos, 2i) = \sin\left(\frac{pos}{10000^{2i/d}}\right)$$

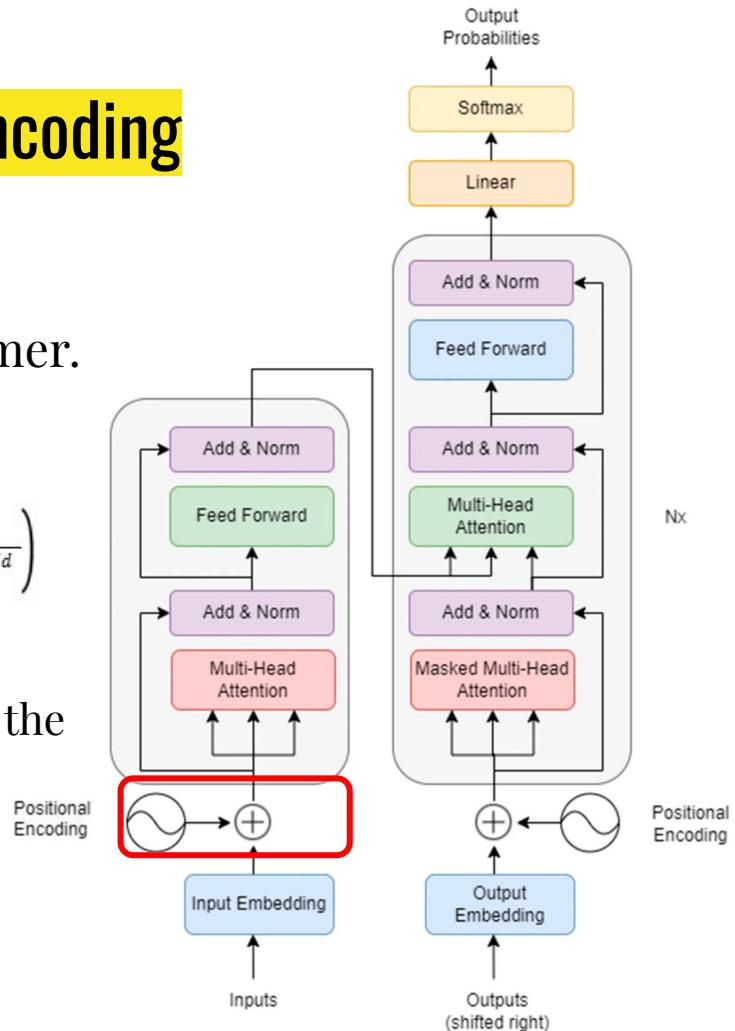
1. pos is the position

$$PE(pos, 2i + 1) = \cos\left(\frac{pos}{10000^{2i/d}}\right)$$

2. i is the embedding dimension

3. d is the dimension of the embedding dimension

sin calculates the even dimensions and cosine calculates the odd dimensions.



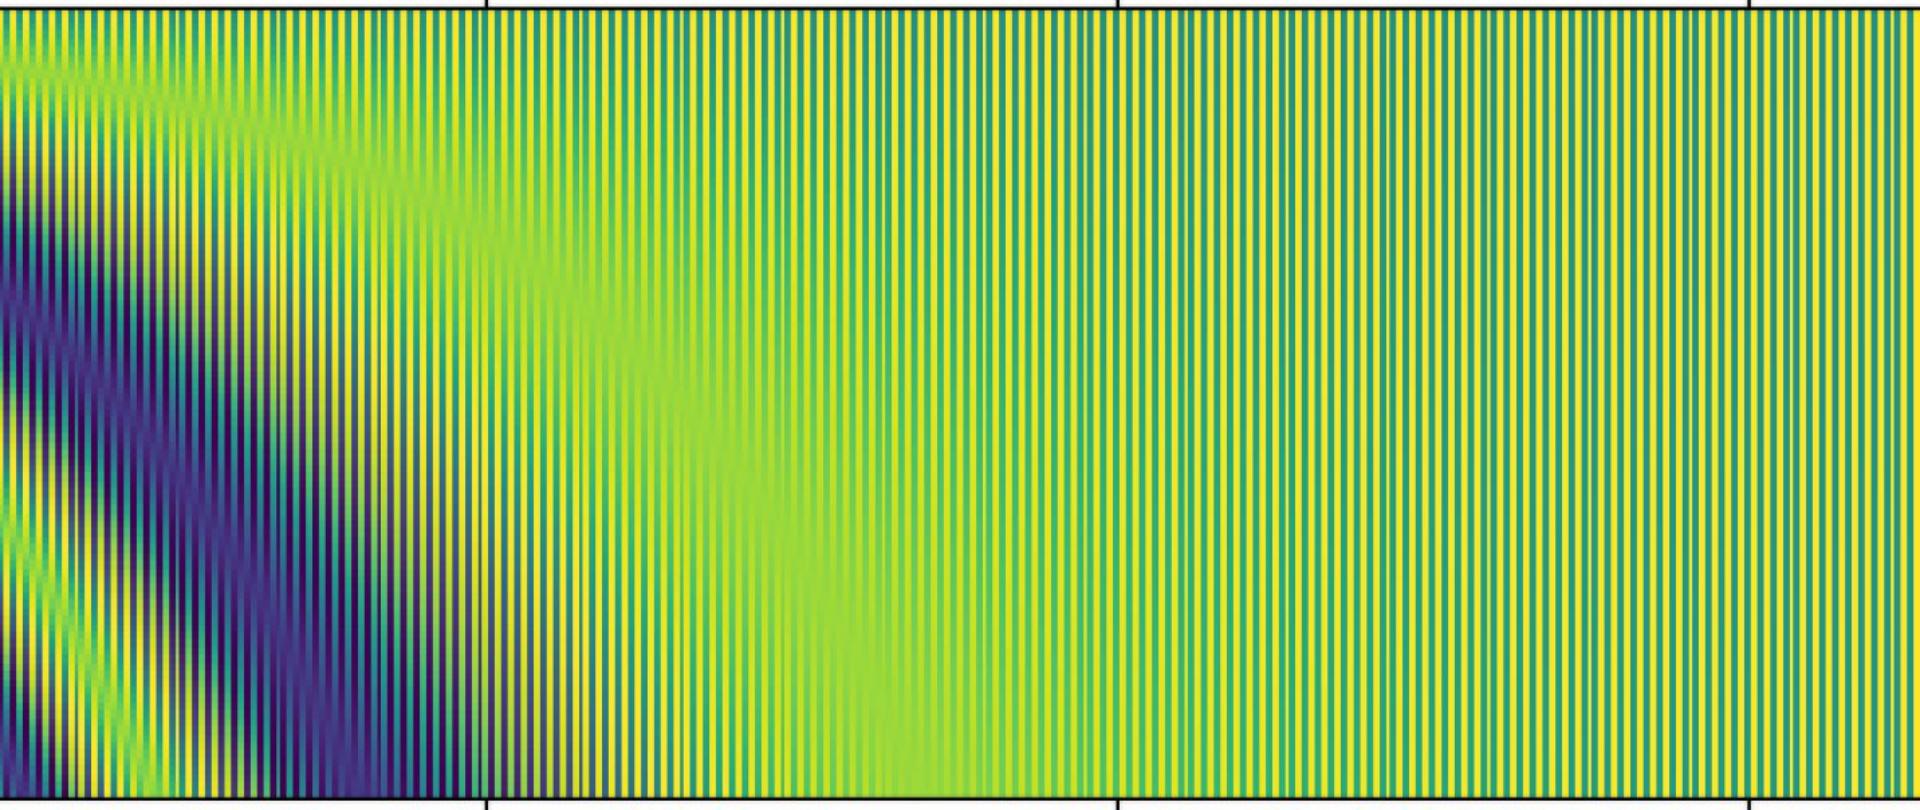
Sinusoidal Positional Encoding Heatmap

200

300

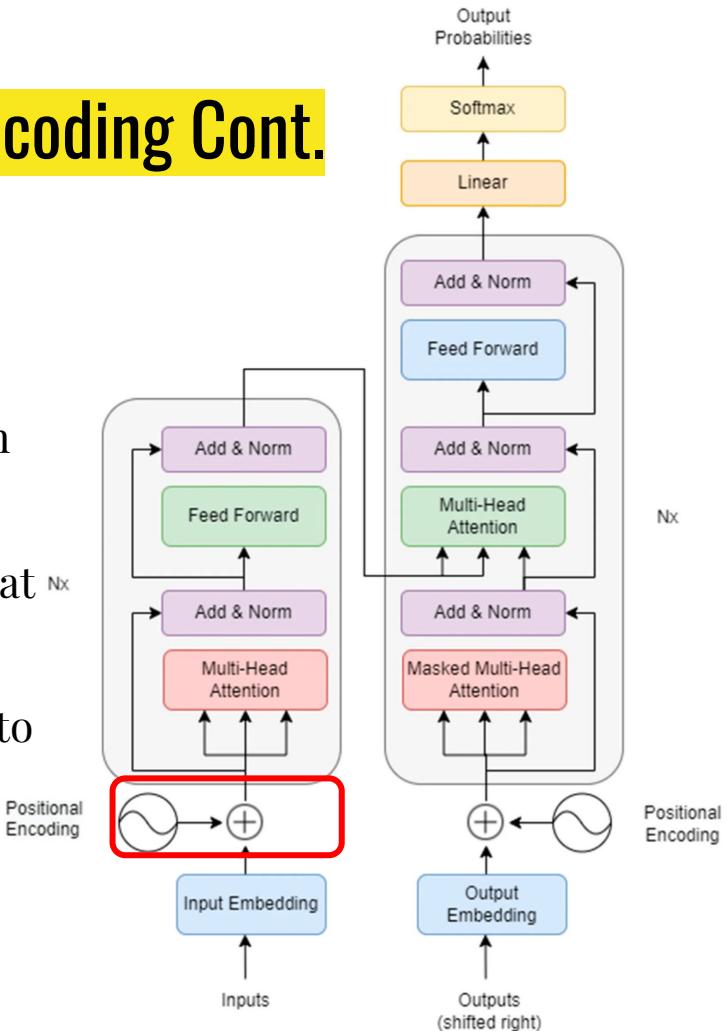
400

Embedding Dimension



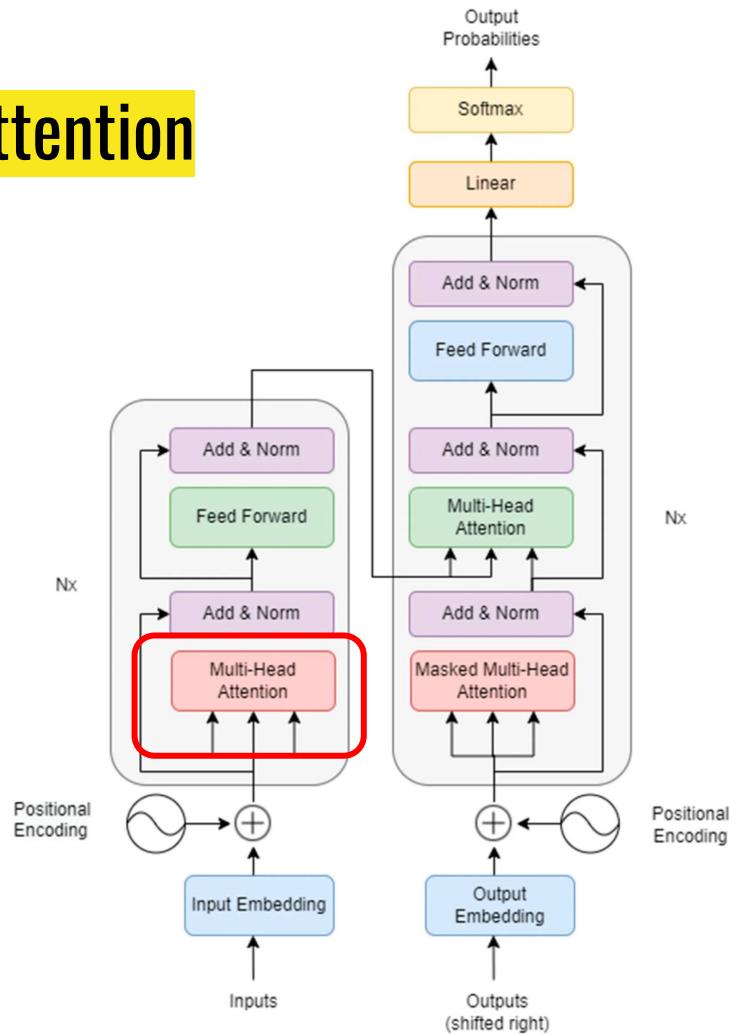
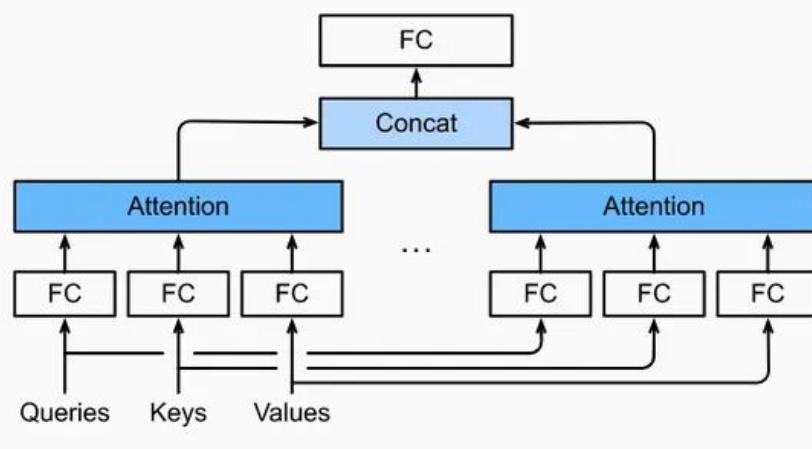
Transformers Architecture - Positional Encoding Cont.

- Learned:
 - ❖ Each position gets a vector embedding. This vector embedding is learned during training such as any other weighted matrix
 - ❖ pos 1 will have a constant embedding regardless of what is at the position.
 - ❖ This is simpler and makes the model not constrained to sinusoidal representation.



Transformers Architecture - Multihead Attention

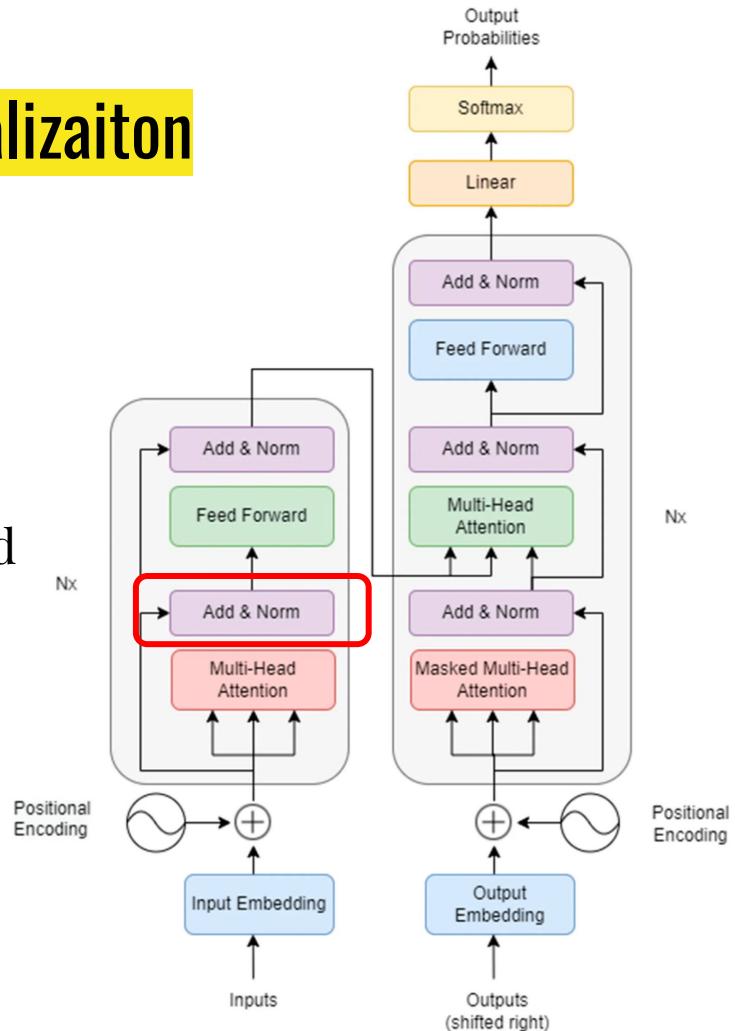
- Instead of relying on one attention head, calculate the attention multiple time and concatenate the results.
- Return to the original dimension by processing through the feed forward layer.



Transformers Architecture - Layer Normalization

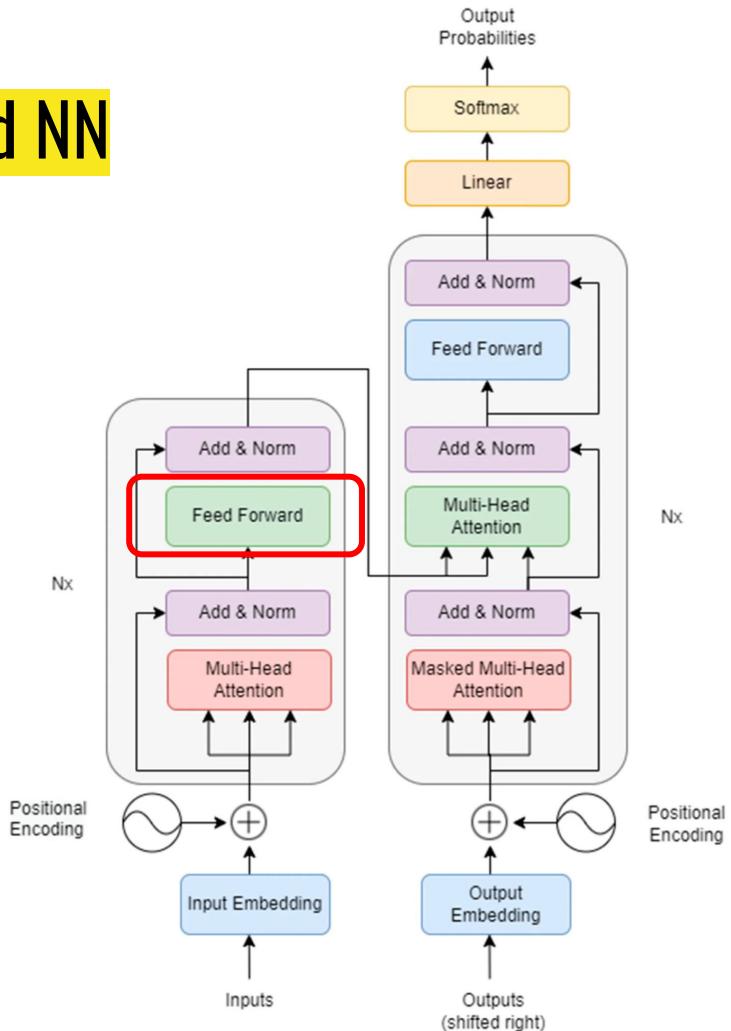
- $x - \mu(x)$

 $\sigma(x)$
- Helps normalize the outputs of the multihead attention layer.



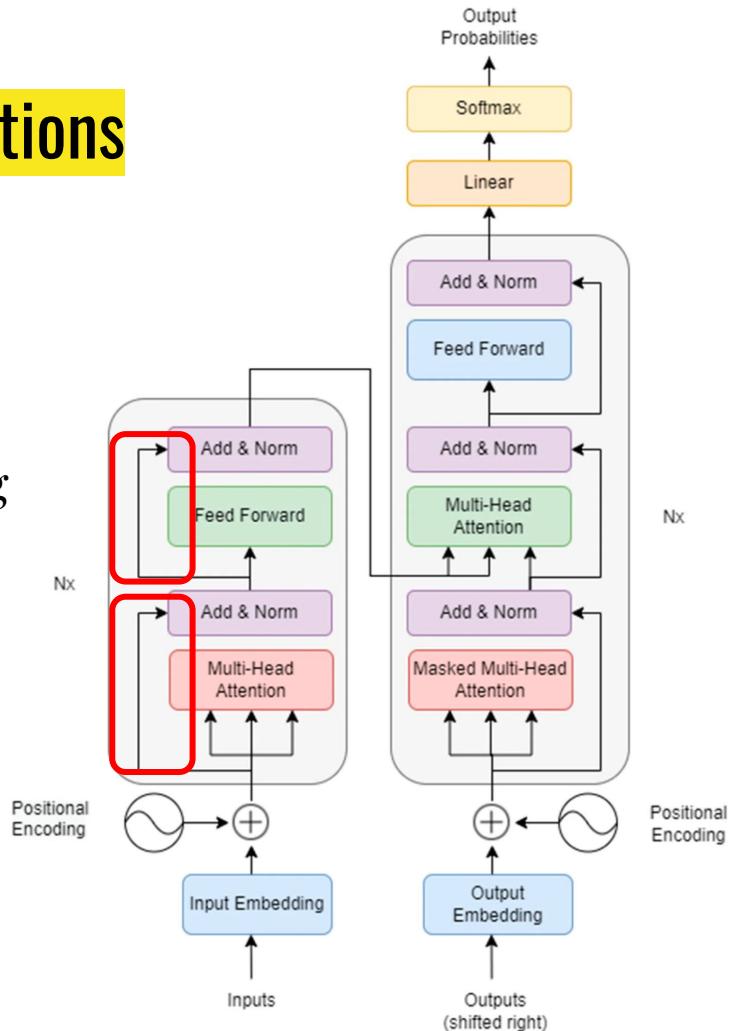
Transformers Architecture - Feed Forward NN

- Feed forward NN introduces the regular NN linear layers with activation functions.
- Preserves the original input shape



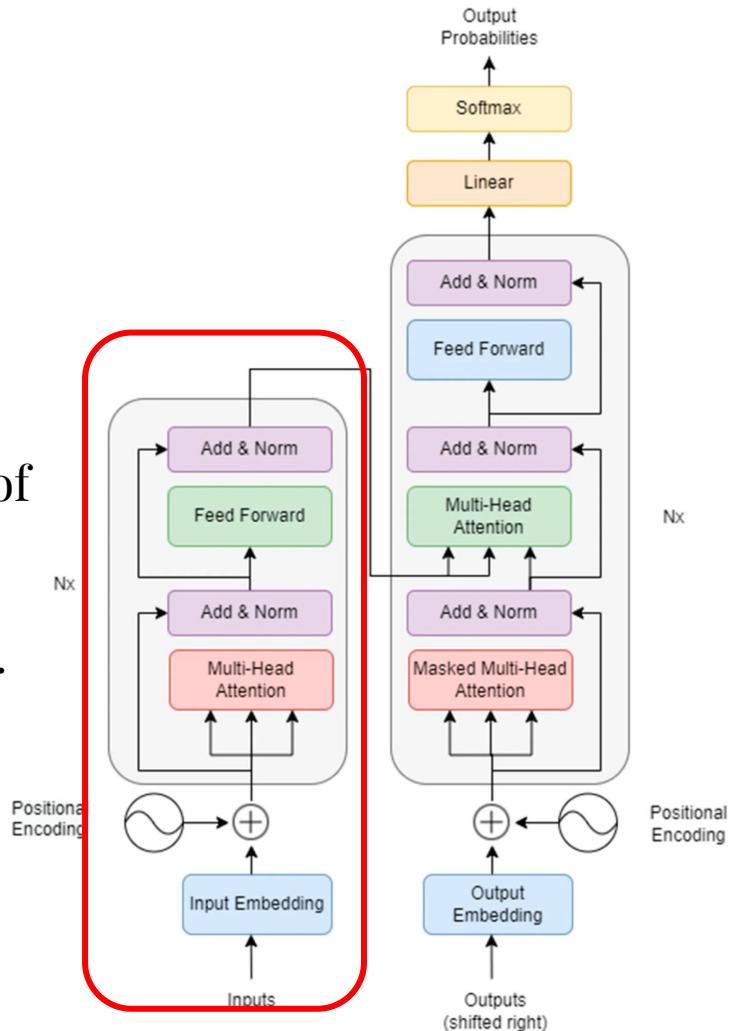
Transformers Architecture - Skip Connections

- As in ResNets, add the output of the previous block to the next block.
- This helps smooth the loss and make training easier.



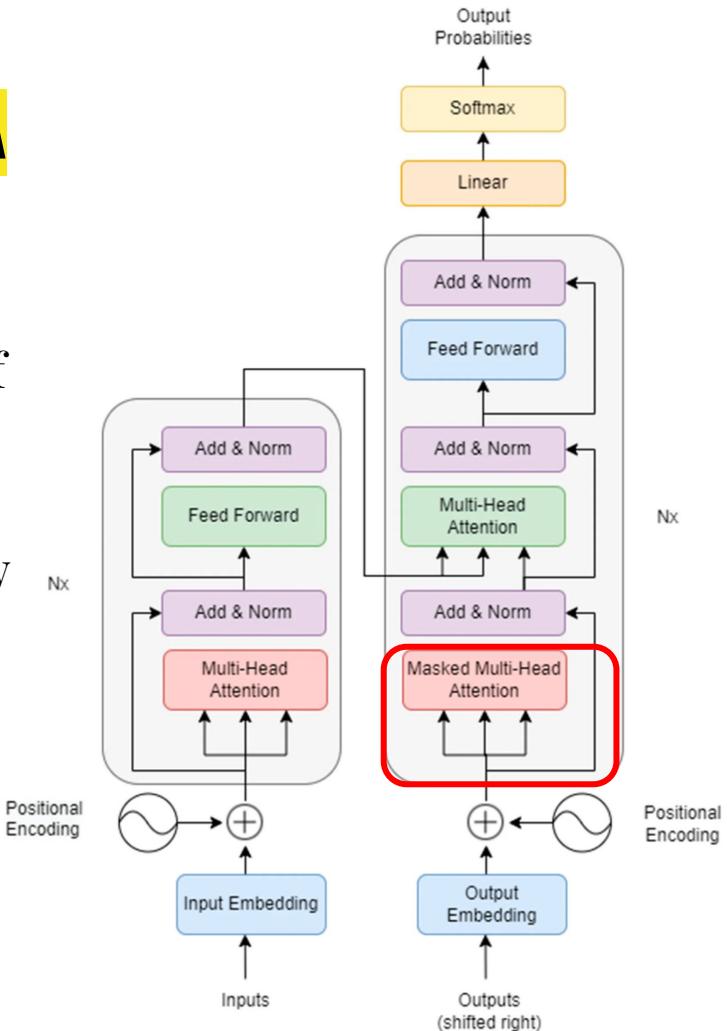
Transformers Architecture - Encoder

- The outputs z can be calculated in parallel which speeds inference using GPUs.
- We can stack more than one encoder on top of each other to introduce complexity to the model and get better contextual embeddings.



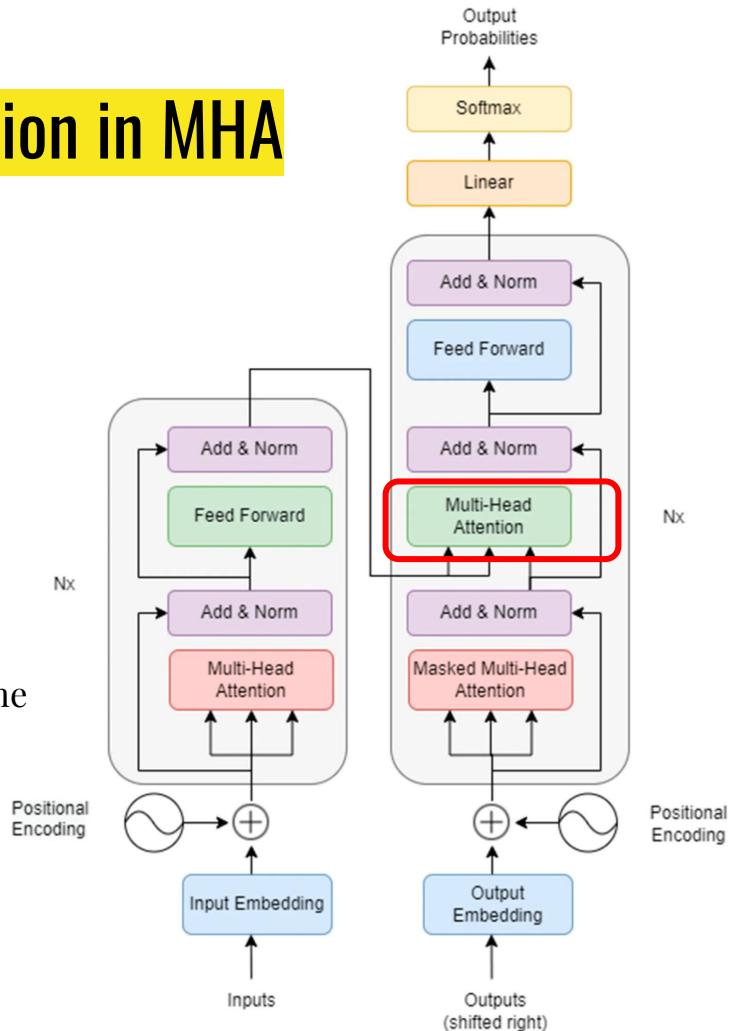
Transformers Architecture - Masked MHA

- Masked multi-head attention is the same as regular attention but “masks” the attention of the outputs.
- “I love dogs.” as the output “I” shouldn’t know that “love” comes after.
- Preserves the “auto regressive” trait.



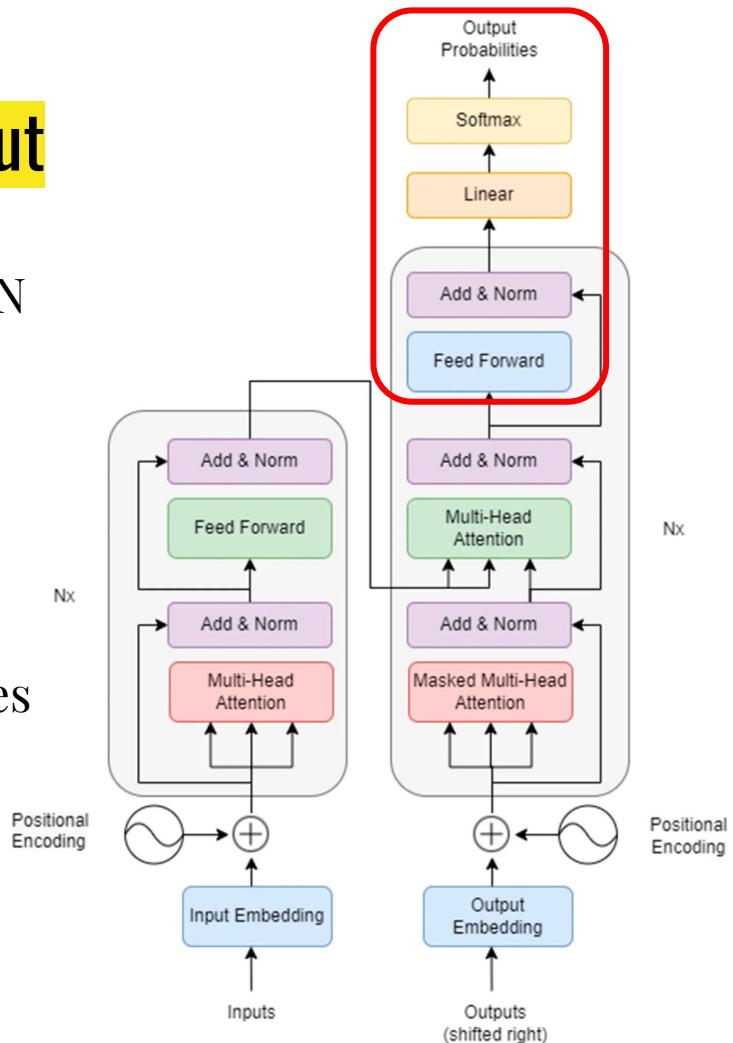
Transformers Architecture - Cross Attention in MHA

- $K = \text{encoder output} \times W_K$ (Key projection matrix)
- $V = \text{encoder output} \times W_V$ (Value projection matrix)
- $Q = \text{decoder hidden} \times W_Q$
- Then it's normal self-attention. This step combines both the encoder and decoder outputs.
- This can be stacked too, but K and V are constant from the same encoder output.



Transformers Architecture - FNN to Output

- The output of MHA goes through a normal NN with linear layers and activation functions.
- This NN and the linear layer that follows reshape the output based on the task.
- The softmax function converts to probabilities in the one-hot encoding space to be able retrieve the sequence in each step.



Code

- Self-Attention
- Transformer from scratch